

Special Problems:

- *Naming*: identify files systemwide
- how are concurrent reads and writes executed?

Naming

Aims:

- *Location Transparency*:
name does not give any hint on location
- *Location independence*:
files can be moved without name being changed

standard approach: *Remote mounting*

Make remote file system available under local name

Achieves only location transparency

Latter difficult to achieve (requires name server)

62

Concurrent Read and Writes

no problem if one central server: can enforce order of operations

With cache, get consistency problems

One possible solution: adopt *session semantics*

Changes visible only after `close`-operation

Problem: File pointers not usable

63

Implementation Issues

Main issue: stateless vs. stateful servers

(Should server keep information about requests?)

Properties of stateless servers

- Fault tolerance
- No `open/close`-requests needed
- No problems with client crashes

Advantage of stateful servers

- Read ahead possible
- Idempotency easier
- File locking possible

Main problem: Cache consistency, especially for stateless servers

64

NFS

Idea: make file systems available on other hosts

Works on different architectures

⇒ Need well-defined protocol

RPC's used for this purpose

Stateless system

⇒ no `open/close` RPC's

Each RPC contains absolute address in file

Caching employed:

- Server does normal caching (no ill-effects)
- Client caches reads and writes
⇒ obtain inconsistency
Data sent to server only when
 - > 8k written
 - file closed on client
 - timeout reached`open` on client checks server

65