

Kernel Architectures

Standard way: *monolithic kernel*:

Only two levels: user mode and kernel mode

All kernel code executed in kernel mode with full privileges

Example: Linux

Communication between parts of OS

Message passing used

Often combined with capabilities for good permission handling

⇒ Efficient message passing vital for performance

Message passing lends itself to asynchronous communication

⇒ bad for implementing Unix system calls

Suitable for embedded systems, in particular special real-time OS

Microkernel

Idea: Restrict amount of code running in kernel mode to minimum

⇒ Implement remainder of OS as services

At bottom: have microkernel with functions like

- Memory Management
- Scheduling
- Low-level device drivers

Higher-level parts like filesystems implemented in user space

Embedded Systems

Chips with power of whole computer systems now in many applications:

- Mobile Phones
- PDA's
- Smart Cards
- On-board controllers of HW

Characterisation of those systems:

- Fewer resources available: memory, storage space
- Often real-time applications necessary (on-board controllers)

Not so much of a problem in general: OS's designed for this case
Only issue: potentially missing MMU
⇒ virtual memory and protection of processes against each other
not implementable
Also paging not available

Have two different kinds of real-time

- 1.) **Hard real-time**: completion required within a guaranteed amount of time
cannot be met by normal time-sharing systems; needs dedicated HW and adaptations to SW
- 2.) **soft real-time**: critical processes receive priority. Requires
 - **pre-emptive priority scheduling** (plus sufficient resources to avoid starvation)
 - **short dispatch latency** (time between arrival of process and start of execution)
Problem: context switch normally only after syscall-completion or when I/O takes place
way out: make **kernel pre-emptible** (eg Solaris 2, newer versions of Linux)
 - **Priority inversion**: increase priority of process if resources required by high-priority process