

Device drivers

Kernel side

View from user space:

Have special file in `/dev` associated with it, together with four systems calls:

- `open`: make device available
- `read`: read from device
- `write`: write to device
- `close`: make device unavailable

Each file may have functions associated with it which are called when corresponding system calls are made

`linux/fs.h` lists all available operations on files

Device driver implements at least functions for `open`, `read`, `write` and `close`.

Automatic recognition of devices

So far: Have seen how devices can be added and used via explicit commands

Nowadays, automatic HW recognition and insertion and removal of devices important

Requires suitable HW support: Each device responds with unique vendor id and product ID when probed

For certain devices (eg usb) device also responds with type (eg usb-storage)

Each device driver keeps a list for which devices and types it is responsible

All device-related information available to user space via `/sys-filesystem`

Special program goes at installation time through all device drivers and records device id's and type

Steps:

- At boot time, kernel probes devices, which respond with unique id indicating vendor and device type
- For each device found, kernel sends info to userspace
- Special program in userspace (`udev`) generates entry in `/dev` and loads appropriate module

Categorising devices

Kernel also keeps track of

- **Physical dependencies** between devices. Example: devices connected to a USB-hub
- **Buses**: Channels between processor and one or more devices. Can be either physical (eg pci, usb), or logical
- **Classes**: Sets of devices of the same type, eg keyboards, mice

Handling Interrupts in Device Drivers

Normal cycle of interrupt handling for devices:

- Device sends interrupt
- CPU selects appropriate interrupt handler
- Interrupt handler processes interrupt
Two important tasks to be done:
 - Data to be transferred to/from device
 - Waking up processes which wait for data transfer to be finished
- Interrupt handler clears interrupt bit of device
Necessary for next interrupt to arrive

Interrupt processing time must be as short as possible

Data transfer fast, rest of processing slow

⇒ Separate interrupt processing in two halves:

- **Top Half** is called directly by interrupt handler
Only transfers data between device and appropriate kernel buffer and schedules software interrupt to start Bottom half
- **Bottom half** still runs in interrupt context and does the rest of the processing (eg working through the protocol stack, and waking up processes)