

# Kernel programming

Kernel has access to **all** resources

Kernel programs not subject to any constraints for memory access or hardware access

⇒ faulty kernel programs can cause system crash

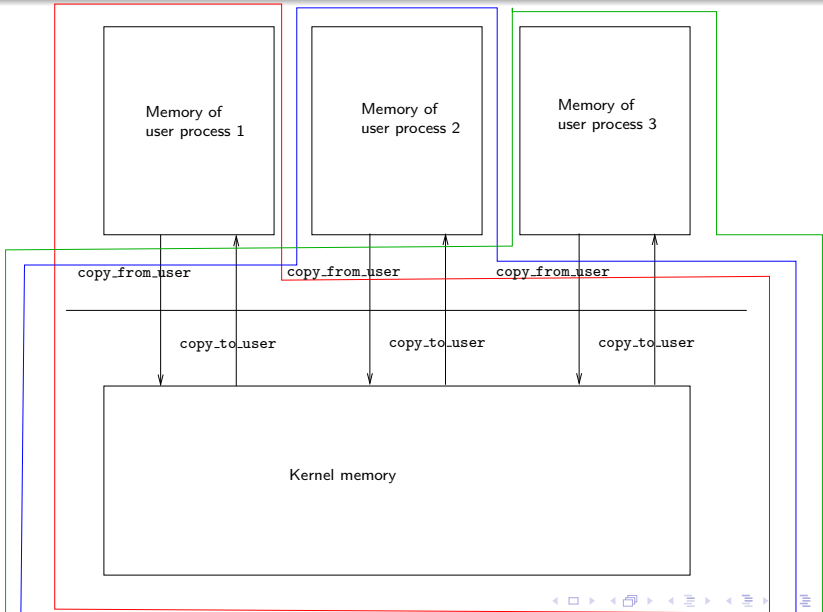
# Interaction between kernel and user programs

Kernel provides its functions only via special functions, called **system calls**

standard C-library provides them

Have strict separation of kernel data and data for user programs

⇒ need explicit copying between user program and kernel



# Structure of kernel

Simplified structure of kernel:

```
initialise data structures at boot time;
while (true) {
    while (timer not gone off) {
        assign CPU to suitable process;
        execute process;
    }
    select next suitable process;
}
```

In addition, have **interrupts**:  
kernel asks HW to perform certain action  
HW sends interrupt to kernel which performs desired action  
interrupts must be processed quickly  
⇒ any code called from interrupts must not sleep

# Linux kernel modes

Structure of kernel gives rise to two main modes for kernel code:

- **process context**: kernel code working for user programs by executing a system call
- **interrupt context**: kernel code handling an interrupt (eg by a device)

have access to user data only in process context

Any code running in process context may be pre-empted at any time by an interrupt

Interrupts have priority levels

Interrupt of lower priority are pre-empted by interrupts of higher priority

# Kernel modules

can add code to running kernel

useful for providing device drivers which are required only if hardware present

`modprobe` inserts module into running kernel

`rmmmod` removes module from running kernel (if unused)

`lsmod` lists currently running modules

# Concurrency issues in the kernel

Correct handling concurrency in the kernel important:  
Manipulation of data structures which are shared between

- code running in process mode and code running in interrupt mode
- code running in interrupt mode

must happen only within critical regions

In multi-processor system even manipulation of data structures shared between code running in process context must happen only within critical sections



# Achieving mutual exclusion

Two ways:

- **Semaphores**: when entering critical section fails, current process is put to sleep until critical region is available  
⇒ only usable if **all** critical regions are in process context
- **Spinlocks**: processor tries repeatedly to enter critical section  
Usable anywhere  
Disadvantage: Have busy waiting

# Programming data transfer between userspace and kernel

Linux maintains a directory called `proc` as interface between user space and kernel

Files in this directory do not exist on disk

Read-and write-operations on these files translated into kernel operations, together with data transfer between user space and kernel

Useful mechanism for information exchange between kernel and user space

# A tour of the Linux kernel

Major parts of the kernel:

- Device drivers: in the subdirectory `drivers`, sorted according to category
- file systems: in the subdirectory `fs`
- scheduling and process management: in the subdirectory `kernel`
- memory management: in the subdirectory `mm`
- networking code: in the subdirectory `net`
- architecture specific low-level code (including assembly code): in the subdirectory `arch`
- include-files: in the subdirectory `include`