


Process Concept

- An operating system executes a variety of programs:
 - Batch system - jobs
 - Time-shared systems - user programs or tasks
- Process - a program in execution; process execution must progress in sequential fashion
- A process includes:
 - program counter
 - stack
 - data section


Process in Memory



Process States

- As a process executes, it changes state
 - new: The process is being created
 - running: Instructions are being executed
 - waiting: The process is waiting for some event to occur
 - ready: The process is waiting to be assigned to a processor
 - terminated: The process has finished execution

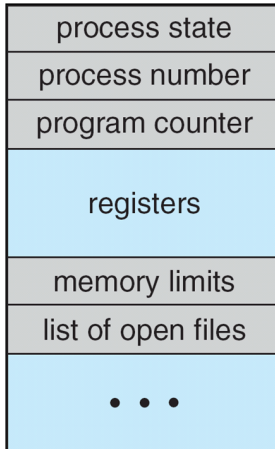
Process States



Process Control Block

- Information associated with each process, which is stored as various fields within a kernel data structure:
 - Process state
 - Program counter
 - CPU registers
 - CPU scheduling information
 - Memory-management information
 - Accounting information
 - I/O status information

Process Control Block




Process Creation

- Parent process create children processes, which, in turn create other processes, forming a tree of processes
- Generally, process identified and managed via a process identifier (pid)
- Resource sharing
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- Execution
 - Parent and children execute concurrently
 - Parent waits until children terminate

Process Creation

- Address space
 - Child duplicate of parent
 - Child has a program loaded into it
- UNIX examples
 - `fork` system call creates new process
 - will look at fork soon, in one of our practical lectures.
 - `exec` system call used after a fork to replace the process' memory space with a new program

Process Creation



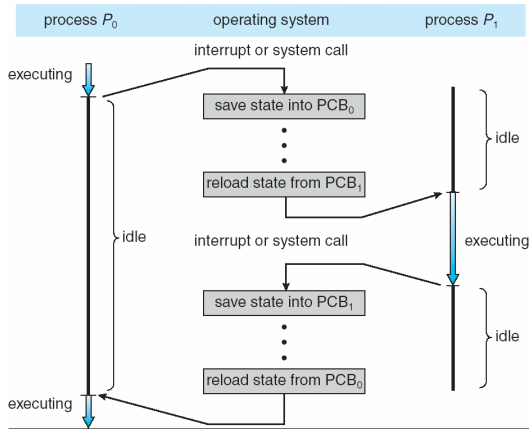
Process Termination

- Process executes last statement and asks the operating system to delete it (exit)
 - Output data from child to parent (via wait)
 - Process' resources are deallocated by operating system
- Parent may terminate execution of children processes (abort)
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - If parent is exiting:
 - Some operating systems do not allow child to continue if its parent terminates — all children terminated (*i.e.* cascading termination)

Context Switch

- When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process via a context switch
- Context of a process represented in the PCB
- Context-switch time is overhead; the system does no useful work while switching
- Time dependent on hardware support

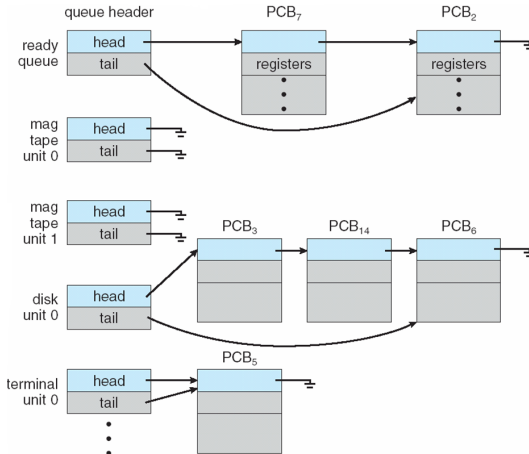
Context Switch



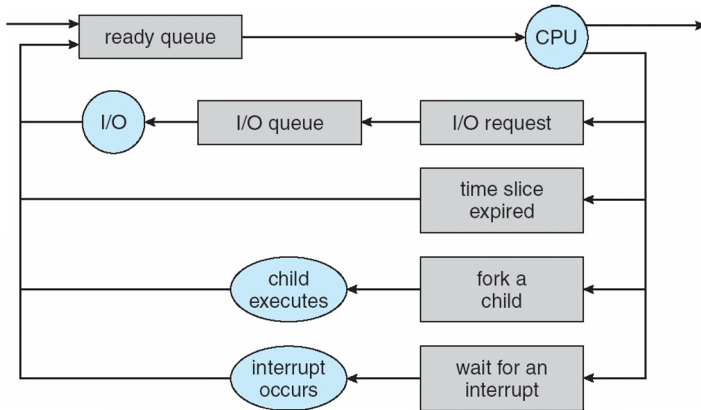
Process Scheduling Queues

- Job queue - set of all processes in the system
- Ready queue - set of all processes residing in main memory, ready and waiting to execute
- Device queues - set of processes waiting for an I/O device
- Processes migrate among the various queues

Process Scheduling Queues



Scheduling Workflow



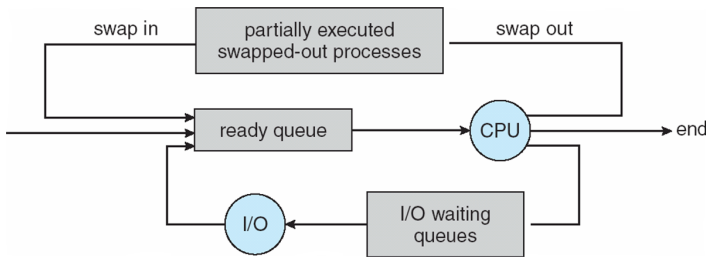
Schedulers

- Long-term scheduler (or job scheduler) - selects which processes should be brought into the ready queue (e.g. loaded from the disk into memory)
- Short-term scheduler (or CPU scheduler) - selects which process should be executed next and allocates CPU

Schedulers

- Short-term scheduler is invoked very frequently (milliseconds) (must be fast)
- Long-term scheduler is invoked very infrequently (seconds, minutes) (may be slow)
- The long-term scheduler controls the degree of multiprogramming (*i.e.* how many processes may compete for the CPU)
 - Long-term scheduling is often minimal or absent on mainstream operating systems, such as Windows and Linux.
- Processes can be described as either:
 - I/O-bound process - spends more time doing I/O than computations, many short CPU bursts
 - CPU-bound process - spends more time doing computations; few very long CPU bursts

Addition of Medium Term Scheduling



- To reduce contention among ready processes, in some scheduling designs, medium-term scheduling allows some processes to be temporarily swapped out of memory
 - kind of like they are being told to sit on a bench until things quieten down.