

Aims:

- *Resource Sharing*: costly resources (high-quality printers and expensive programs) can be shared
- *Computation Speedup*: Algorithms can run concurrently
- *Reliability*: Failure of one site does not imply failure of the whole system
Redundancy prerequisite

Especially last point difficult to fulfill

23

Several Levels of Distribution possible
listed from tightly coupled to loosely coupled

- Shared memory
- Shared file system
- Bus Systems
- Switching Systems

Key Problem Areas:

- Transparency: pretend to be one computer
- Reliability: want availability, fault tolerance
- Performance
- Scalability: avoid centralised tables and algorithms

24

Communication

Simplest way of communication:
Client/Server Model

Idea: group of processes (servers) used by clients

Advantage: Simple communication

Simple enough to study several problems

First problem: Addressing! Possible Solutions:

- Hardware address into client code: inflexible
- Broadcasting: works only on local networks
- Name Servers: Ask special host
Example: Domain Name Service, DNS

25

Second Problem:

Blocking (synchronous) vs. non-blocking primitives

Conceptual ease vs. performance

Third Problem:

Reliable vs. unreliable primitives

Where does the error correction go?

Kernel (once for all) vs. application (possibly more efficient)

26

Very simple idea: execute procedure on different host

Goal: total transparency

Basic Schema:

- Client sends arguments to server
- Server executes call
- Server sends results back

Difference to local call hidden in kernel routines

Details complicated:

- Have to transfer parameters
- deal with failures

27

Problems with parameter passing

- Different representation on different machines: either common format (inefficient), or store format in message
- What to do with call-by-reference parameters? Can copy arrays, but not arbitrary data structures

28

Failure problems more complicated
Have several cases

- Client cannot locate server: Generate exception
⇒ transparency lost
- Lost Request Messages: use timer
- Lost Reply Messages: client timer insufficient:
could execute operation more than once
Solution: use sequence numbers
- Server Crashes: Sequence numbers not enough: When did crash occur?
Can guarantee *at least once*, at *most once* semantics, but not *exactly once* semantics ⇒ have to have call-specific remedies
- Client Crash: leaves orphans (unwanted computations)
No general way of getting rid of them

29

Generating Timestamps

Unique timestamps needed for co-ordination
No problem in monoprocessor system: use system clock

Not possible in distributed systems

One way out:

- Each host maintains logical clock which is advanced with each event
- All message from host contain logical clock
- When message with greater logical clock is received, increase own logical clock to this value
- with two identical timestamps, let host number decide

30

A Distributed Version (Ricart and Agrawala)

Assume reliable messages, unique timestamps

Following steps:

- Process trying to enter critical section:
sends to all other processes name of section and unique timestamp
- Process receiving such a message:
 - Sends back OK if not interested in critical section
 - Queues message if already in critical section
- Receiver wants to enter critical section
⇒ enters critical section if its request has lower timestamp and queues message, otherwise sends OK

31

Grants mutual exclusion without deadlock or starvation

Problems:

- Requires that everyone knows about all other hosts
- Algorithm fails if one host fails
- Load for *all* hosts higher than for coordinator in centralised algorithm

32

A Token Ring Algorithm

Assumption: Network organised on a (physical or logical) ring, *i.e.*,

each node has unique successor in line

Simple algorithm:

- At initialisation, generate token
- Pass token around continuously
- Process wanting to enter a critical region waits for token
- After leaving critical section, process passes token to next neighbour

Properties:

- Detecting lost tokens difficult: Time spent in critical region unbound
- Detecting dead processes easy if sent token is acknowledged

33

Comparison of three algorithms:

- Distributed algorithms *more* sensitive to crashes than centralised one
- Central algorithms simplest and most efficient

⇒ Distributed systems work only with reliable components

34

Problem: Select new co-ordinator

Assumption:

Know id of every host on the network

First example: Bully algorithm

- P sends message to all hosts with higher number
- No response $\Rightarrow P$ wins and is co-ordinator
- Answer received \Rightarrow host with higher number has taken over

Second Example: Ring Algorithm

- any host may send *Election* message
- passed around the net, which each hostid added
- If original host gets message, determines co-ordinator and sends new message around
- After it has gone round, host removes it