

OS accessible only via specified procedures
(system calls)

Execution of a system call:

- OS reads call parameters and checks appropriate privileges
- OS executes requested function in Supervisor mode
- OS returns result

11

Informally: Process is a program in execution.

States of a process:

New Process is being created


Running Instructions are being executed

Waiting Process waiting for an event to occur
(I/O; signal)

Ready Process waiting for processor only

Terminated Process has finished execution

12



13

Process Control Block

OS allocates one Process Control Block (PCB) per process to store all relevant information, e.g.,

- Process state
- Program counter
- CPU-registers
- CPU-scheduling information
- Memory management information
- Accounting information
- I/O-status information

PCB used to identify process

14

Ready-state associated with a **queue of processes** waiting to be executed
Scheduler picks one from the queue according to some criteria (later)

Instance of **general principle**:

Whenever we have a shared resource (here: processor), **collect all waiting processes in a queue** and use scheduler to decide who is next

Changing running process requires **saving state of old one and loading state of new one**
⇒ **significant** overhead involved (upto 1 ms).

Amount of work increases with complexity of OS

⇒ has become a **real performance bottleneck**

Will discuss one possible remedy (**Threads**) later

Special **system call** for this (`fork` in UNIX)

Creating process is called **parent**, created process **child**.

Several possibilities for **relation between parent and child**:

- Concurrent execution vs. parent waiting for children
- Child is duplicate of parent or has separate code

Often all **variants can be realised** (UNIX).

Problem: **Context switch takes too long**
⇒ Can we reduce amount of material to be stored and re-fetched?

Solution: differentiate between

- **data**: memory, file pointers ...
- **processor information**: program counter, register set, stack

Processes manage all these data
introduce **Threads**, which are **activity carriers together with processor information**

Two possibilities for implementation of threads:

- **User-level:** OS doesn't know about them, only have a library. \Rightarrow Minimal overhead, but one blocked thread of process blocks whole process
- **kernel-level:** implementation done in OS: slower, but fairer allocation of resources possible

Some OS (e.g. Solaris) provide both versions.

Threads require good handling of concurrent processes.